

Image Processing and Computer Graphics

Image Processing

Class 4 Energy Minimization

1. Formalize your model assumptions and cast the task as an **optimization problem**

$$E(x) = A_1(x) + \dots + A_n(x)$$

2. Solve the optimization problem

$$x^* = \operatorname{argmin}_x E(x)$$

- Objective function $E(x)$ often called **energy** (motivated from physics)
- In machine learning it is often called **loss function**

Example: image denoising

- First step: formulate the model assumptions
 - The outcome should be similar to the input image
 - The result should be smooth

- Second step: formalize these assumptions

- Similarity to the input data (data term):

$$E_D(u_{i,j}) := \sum_{i,j} (u_{i,j} - I_{i,j})^2 \rightarrow \min$$


- Similarity to neighboring values (smoothness term):

$$E_S(u_{i,j}) := \sum_{i,j} (u_{i+1,j} - u_{i,j})^2 + (u_{i,j+1} - u_{i,j})^2 \rightarrow \min$$

- Yields an energy minimization problem including a weighting parameter α

$$u_{i,j}^* = \operatorname{argmin}_{u_{i,j}} (E_D(u_{i,j}) + \alpha E_S(u_{i,j}))$$

- Third step: solve this optimization problem



- All model assumptions are clearly stated
→ transparency
- Global approach: all variables are optimized in a joint manner;
interdependencies are not lost by intermediate decisions → optimality
- Theoretical aspects can be analyzed:
 - Existence and uniqueness of solutions
 - Stability of solutions with respect of the input data
 - Difficulty of the problem class
- Usually fewer parameters than in heuristic multi-step methods
- Combination of energy minimization approaches is straightforward

- Formalizing the assumptions is rather ad-hoc: Why minimizing the sum of squared differences and not some other distance?

$$E_D(u_{i,j}) := \sum_{i,j} (u_{i,j} - I_{i,j})^2 \rightarrow \min$$

→ Statistical interpretations (Bayesian models) can usually answer this question.

- Choosing the weighting parameter(s) is not easy and often depends on the data.
→ (Fixed) parameters can be learned from a training dataset.
- Global optimization is often hard.
Heuristics obliterate the initial transparency of the model.

- Here is our energy from the denoising example

$$E(u_{i,j}) := \sum_{i,j} \left((u_{i,j} - I_{i,j})^2 + \alpha \left((u_{i+1,j} - u_{i,j})^2 + (u_{i,j+1} - u_{i,j})^2 \right) \right)$$

- How do we find the minimum?
- Necessary condition** for a minimum: the first derivatives must be zero

$$\frac{dE}{du} = 0 \quad \Leftrightarrow \quad \frac{\partial E}{\partial u_{i,j}} = 0 \quad \forall i, j$$

- Here we go:

$$\begin{aligned} \frac{\partial E}{\partial u_{i,j}} = & 2(u_{i,j} - I_{i,j}) \\ & + 2\alpha(u_{i,j} - u_{i-1,j}) - 2\alpha(u_{i+1,j} - u_{i,j}) \\ & + 2\alpha(u_{i,j} - u_{i,j-1}) - 2\alpha(u_{i,j+1} - u_{i,j}) = 0 \end{aligned}$$

- At boundary pixels some terms are missing due to missing neighbors

- Necessary conditions...

$$\begin{aligned} \frac{\partial E}{\partial u_{i,j}} = & (u_{i,j} - I_{i,j}) \\ & + \alpha(u_{i,j} - u_{i-1,j}) - \alpha(u_{i+1,j} - u_{i,j}) \\ & + \alpha(u_{i,j} - u_{i,j-1}) - \alpha(u_{i,j+1} - u_{i,j}) = 0 \end{aligned}$$

- ...can be written as a large linear system of equations (schematic view)

$$\begin{pmatrix} 1+2\alpha & -\alpha & & -\alpha & & & \\ -\alpha & 1+3\alpha & -\alpha & & -\alpha & & \\ & -\alpha & 1+3\alpha & -\alpha & & -\alpha & \\ -\alpha & & -\alpha & 1+4\alpha & -\alpha & & -\alpha \\ & -\alpha & & -\alpha & 1+3\alpha & -\alpha & \\ & & -\alpha & & -\alpha & 1+3\alpha & -\alpha \\ & & & -\alpha & & -\alpha & 1+2\alpha \end{pmatrix} \begin{pmatrix} u_1 \\ u_2 \\ \vdots \\ \vdots \\ u_{N-1} \\ u_N \end{pmatrix} = \begin{pmatrix} I_1 \\ I_2 \\ \vdots \\ \vdots \\ I_{N-1} \\ I_N \end{pmatrix}$$

- $N \times N$ system matrix (for N pixels) is symmetric and positive definite
- It contains one main diagonal (central pixels) and four off-diagonals (for each of the four neighbors of a pixel)


$$\begin{pmatrix} 1+2\alpha & -\alpha & & & & \\ -\alpha & 1+3\alpha & -\alpha & & & \\ & -\alpha & 1+3\alpha & -\alpha & & \\ -\alpha & & -\alpha & 1+4\alpha & -\alpha & \\ & -\alpha & & -\alpha & 1+3\alpha & -\alpha \\ & & -\alpha & & -\alpha & 1+3\alpha \\ & & & -\alpha & & -\alpha & 1+2\alpha \end{pmatrix} \begin{pmatrix} u_1 \\ u_2 \\ \vdots \\ \vdots \\ u_{N-1} \\ u_N \end{pmatrix} = \begin{pmatrix} I_1 \\ I_2 \\ \vdots \\ \vdots \\ I_{N-1} \\ I_N \end{pmatrix}$$

- The system matrix A is **sparse** (almost all entries are 0)
- Positive definite \rightarrow the inverse A^{-1} exists and we can solve for \mathbf{u}
- Questions:
 - When do we get such a linear system with a unique solution?
 - How can this system be solved (efficiently)?

- Convex functions:
 - Positive curvature
 - No local minima
 - Global minimum is unique
 - Minimization by setting the derivative to 0 and solving the emerging linear or nonlinear system



- Non-convex functions:
 - Usually many local minima (and maxima)
 - Global minimum may not be unique
 - Global minimization is usually impossible, only heuristics exist




- A function is **convex** if

$$f((1-\alpha)x_1 + \alpha x_2) \leq (1-\alpha)f(x_1) + \alpha f(x_2) \quad \forall x_1, x_2, \forall \alpha \in (0, 1)$$

- A function is **strictly convex** if

$$f((1-\alpha)x_1 + \alpha x_2) < (1-\alpha)f(x_1) + \alpha f(x_2) \quad \forall x_1, x_2, \forall \alpha \in (0, 1)$$



- Theorem: every convex combination of (strictly) convex functions is again (strictly) convex

- Proof:

$$h(x) := \gamma f(x) + \delta g(x) \quad \gamma, \delta \in \mathbb{R}, \quad \gamma, \delta \geq 0$$

$$\begin{aligned} h((1 - \alpha)x_1 + \alpha x_2) &= \gamma f((1 - \alpha)x_1 + \alpha x_2) + \delta g((1 - \alpha)x_1 + \alpha x_2) \\ &\leq \gamma(1 - \alpha)f(x_1) + \gamma\alpha f(x_2) + \delta(1 - \alpha)g(x_1) + \delta\alpha g(x_2) \\ &= (1 - \alpha)(\gamma f(x_1) + \delta g(x_1)) + \alpha(\gamma f(x_2) + \delta g(x_2)) \\ &= (1 - \alpha)h(x_1) + \alpha h(x_2) \end{aligned}$$

- Our energy function

$$E(u_{i,j}) := \sum_{i,j} \left((u_{i,j} - I_{i,j})^2 + \alpha \left((u_{i+1,j} - u_{i,j})^2 + (u_{i,j+1} - u_{i,j})^2 \right) \right)$$

is a convex combination of strictly convex (quadratic) functions

→ It is strictly convex as well

→ It has a unique global minimum

$$\begin{pmatrix}
 1+2\alpha & -\alpha & & & \\
 -\alpha & 1+3\alpha & -\alpha & & \\
 & -\alpha & 1+3\alpha & -\alpha & \\
 -\alpha & & -\alpha & 1+4\alpha & -\alpha \\
 & -\alpha & & -\alpha & 1+3\alpha \\
 & & -\alpha & & -\alpha & 1+3\alpha & -\alpha \\
 & & & -\alpha & & -\alpha & 1+2\alpha
 \end{pmatrix}
 \begin{pmatrix}
 u_1 \\
 u_2 \\
 \vdots \\
 \vdots \\
 u_{N-2} \\
 u_{N-1}
 \end{pmatrix}
 =
 \begin{pmatrix}
 I_1 \\
 I_2 \\
 \vdots \\
 \vdots \\
 I_{N-2} \\
 I_{N-1}
 \end{pmatrix}$$

- The two additional off-diagonals (together with the size of matrix) rule out Gauß-elimination (in 1D, however, Gauß-elimination is very efficient).
- An iterative solver is needed to preserve the sparsity of the matrix
- Simplest iterative solver: **Jacobi method**
- Converges if the matrix is **strictly diagonal dominant**

$$|a_{ii}| > \sum_{i \neq j} |a_{ij}| \quad \forall i$$

- Decompose the matrix into its diagonal part D and its off-diagonal part M
 $A = D + M$

- For the linear system this means

$$Ax = b \quad \Leftrightarrow \quad (D+M)x = b \quad \Leftrightarrow \quad Dx = b - Mx$$

- D^{-1} can be computed very easily: just replace the diagonal elements by their inverse.
- Now we can compute the solution x iteratively. Starting with any initialization x^0 , iterate

$$x^{k+1} = D^{-1}(b - Mx^k)$$

- Iterate until the norm of the **residual** $r^k := Ax^k - b$ is smaller than a threshold or the change in the solution $(x^{k+1} - x^k)^2$ becomes small. When the change is 0, the iterate has **converged**.

- Advantages:
 - Simple
 - Can be implemented in parallel
- Disadvantages:
 - Slow
 - Convergence only for $k \rightarrow \infty$
 - Computation not in-place
- Alternatives:
 - Gauß-Seidel, Successive Over-relaxation (faster, in-place)
 - Conjugate gradient (convergence after finite number of iterations)
 - Multigrid methods (sometimes much faster)

- Split the off-diagonal part M into the lower triangle L and the upper triangle U

$$Ax = b \quad \Leftrightarrow \quad Dx = b - Lx - Ux$$

- During iteration, traverse the vector x from top to bottom and use already the new values for multiplication with the lower triangle

$$x^{k+1} = D^{-1}(b - Lx^{k+1} - Ux^k)$$

- In our denoising example this reads

$$u_i^{k+1} = \frac{I_i + \alpha \sum_{j \in \mathcal{N}^-(i)} u_j^{k+1} + \alpha \sum_{j \in \mathcal{N}^+(i)} u_j^k}{1 + \sum_{j \in \mathcal{N}(i)} \alpha}$$

- Converges if A is positive or negative definite
- For already updated neighbors take the new value \rightarrow in-place computation
- Recursive propagation of information \rightarrow faster

- Emphasize the Gauß-Seidel idea by over-relaxing the new solution

$$x^{k+1} = (1 - \omega)x^k + \omega D^{-1}(b - Lx^{k+1} - Ux^k)$$

- For $\omega = 1$ this is the Gauß-Seidel method
- Converges for positive- or negative-definite matrices (all eigenvalues positive or negative, respectively), if $\omega \in (0, 2)$
- Over-relaxation for $\omega > 1$: faster convergence
- Under-relaxation for $\omega < 1$: can help establish convergence in case of divergent iterative processes
- Optimal ω must be determined empirically

- Two non-zero vectors u and v are **conjugate** with respect to A if the inner product $\langle u, v \rangle_A := u^\top A v = 0$. This means the two vectors are orthogonal with respect to this special scalar product.
- A set of n conjugate vectors $\{p_k\}$ forms a basis of \mathbb{R}^n , so the solution x^* of $Ax = b$ can be expanded as $x^* = \alpha_1 p_1 + \dots + \alpha_n p_n$
- The coefficients α_k are derived as follows:

$$Ax^* = \alpha_1 A p_1 + \dots + \alpha_n A p_n = b$$

$$p_k^\top A x^* = p_k^\top \alpha_1 A p_1 + \dots + p_k^\top \alpha_n A p_n = p_k^\top b \quad (\text{expansion with } p_k)$$

$$\alpha_k = \frac{p_k^\top b}{p_k^\top A p_k}$$

- After n computations we obtain the exact solution x^*
- Good choice of $\{p_k\} \rightarrow$ few coefficients approximate the solution well

- Start with some initial point x^0
- Let p_0 be the residual $r_0 = b - Ax^0$. This is the gradient of

$$E(x) = \frac{1}{2}x^\top Ax - b^\top x$$

the minimizer of which is x^* . Therefore the name conjugate gradient.

- Iteratively compute:

$$\alpha_k = \frac{r_k^\top r_k}{p_k^\top A p_k} \quad x^{k+1} = x^k + \alpha_k p_k \quad r_{k+1} = r_k - \alpha_k A p_k$$

$$\beta_k = \frac{r_{k+1}^\top r_{k+1}}{r_k^\top r_k} \quad p_{k+1} = r_{k+1} + \beta_k p_k$$

- Stop when residual is small. Guaranteed solution after n iterations.

- Matrix A must be symmetric and positive definite
- Usually in image processing, computing the exact solution is not an option since n is the number of pixels
- Number of iterations needed to get a good approximate solution depends on the **condition number** of A (largest vs. smallest eigenvalue). The same holds for the other iterative methods (Gauß-Seidel, etc.)
- Sometimes so-called **preconditioners** P^{-1} are used to have a small condition number for $P^{-1}A$

$$Ax = b \quad \Leftrightarrow \quad P^{-1}Ax = P^{-1}b$$

- Simplest preconditioner: Jacobi preconditioner



$$P = D \quad \Leftrightarrow \quad P^{-1} = D^{-1}$$

$$Ax = b \quad \Leftrightarrow \quad D^{-1}Ax = D^{-1}b$$

- All previous linear solvers have the drawback that they only act locally.
- This is due to the sparsity of the matrix: one iteration only distributes information at a pixel to its four neighbors (the recursive nature of Gauß-Seidel relaxes this statement a bit).
- Idea of multigrid solvers: shorten distances by regarding the system from a coarser point of view
- Additional effect: coarse versions of the system have fewer entries
→ iterations are faster at coarse levels
- Different types of multigrid solvers
 - Unidirectional (cascadic) multigrid
 - Bidirectional (correcting) multigrid
 - Full multigrid (a combination of both)

- Create downsampled versions of the linear system (usually by downsampling the image and deriving the linear system from this)
- Compute first approximate solution at the coarse grid (e.g. with SOR)
- Take upsampled result as initial guess for the next finer grid
- Refine result there (again with SOR)
- Advantages:
 - Iterations at the coarse level are very fast
 - Simple implementation
- Disadvantage:
 - Coarse level systems often do not approximate the original system well
→ Iterations at coarse levels do not really help


- Basic idea: do not downsample the image but the error
- Compute first solution at fine grid
- Correct error at coarse grid
- Refine result at finer grid



Author: Andrés Bruhn

1. Presmoothing relaxation step $A^h x^h = b^h$

- Run some iterations at the fine grid
- Yields approximate solution \tilde{x}^h
- Remaining error: $e^h = x^h - \tilde{x}^h$




2. Correction step:

- Goal: compute error $A^h e^h = r^h$ $r^h = b^h - A^h \tilde{x}^h$
- Local part of error already removed
→ Solve this system at coarser grid $A^H e^H = r^H$
- Transfer error to fine grid and correct the solution $\tilde{\tilde{x}}^h = \tilde{x}^h + \tilde{e}^h$

3. Postsmoothing relaxation step

- Apply some further iterations at fine grid to remove local errors introduced by \tilde{e}^h

- Combination of cascadic and correcting multigrid
- Start at coarse grid with downsampled image
- At each finer level apply a W-cycle



Author: Andrés Bruhn

- Here we were considering problems with continuous variables (each vector component of the solution is a real number)
- Segmentation and matching typically leads to integer problems
- These are combinatorial problems, only few of them being solvable in polynomial time
- Some typical problems arising in computer vision are:
 - Linear programs (LP)
 - Integer quadratic programs (IQP) including special cases like min-cut
 - Second order cone programs (SOCP)
- More in the Computer Vision course

- The energy minimization framework is a sound way to model and solve image processing problems
- All model assumptions are clearly stated, no hidden assumptions
- Global optimization is “easy” if the energy function is convex
- The necessary condition for a minimum is that the gradient is zero
- Leads to a large, but sparse, linear or nonlinear system of equations
- There are several methods to solve sparse linear systems iteratively, some are easier to implement, others are faster

- D. Young: Iterative Solution of Large Linear Systems, Academic Press, 1971. Reprint by Dover 2003.
- J. R. Shewchuk: An introduction to the Conjugate Gradient method without the agonizing pain. CMU Technical Report, 1994.