

Image Processing and Computer Graphics

Image Processing

Class 8

Interest points and local descriptors

- Key problem in computer vision appearing in:
 - Motion estimation
 - Camera calibration
 - Stereo
 - Image retrieval
 - Object recognition




Object recognition: training image on the left, test image on the right. Matching here is quite hard.




Stereo pair: point matching needed to compute depth

- Straightforward way to match points in images:
 - Regard the image patch around each point in image 1
 - Compare it to the image patches around all points in image 2



- Computationally expensive
 $O(kN^2)$, k : size of patch, N : size of image (in pixels)
- Not **invariant** to typical appearance changes

- Often we need only a limited number of matches
- Idea: Do not match all points in the images, but only promising subsets
→ significantly reduced complexity
- Requirements for good interest points:
 1. Points must come with enough information for unique matching



2. Subset in image 2 must contain matches from subset in image 1

- Choose points with high information content and clear localization
→ typically corner points
- Corner detection with the structure tensor:
(Förstner-Gülch 1987, Harris-Stevens 1988)

$$J_\rho = K_\rho * (\nabla I \nabla I^\top) = \begin{pmatrix} K_\rho * I_x^2 & K_\rho * I_x I_y \\ K_\rho * I_x I_y & K_\rho * I_y^2 \end{pmatrix}$$

- Measure of cornerness (fast to compute):


$$c = \det J_\rho - \alpha \operatorname{tr} J_\rho$$

\uparrow = gradient magnitude

- Eigenvalue decomposition of the structure tensor:

$$J_\rho = T \begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix} T^\top$$

$$c = \lambda_2$$



Input and second eigenvalue for different ρ

- Interpretation:
 - Smoothing of J integrates gradients from the neighborhood
 - Eigenvectors in T yield the dominant orientation in this neighborhood and the perpendicular orientation
 - Eigenvalues yield the structure magnitude in these directions
 - A large second eigenvalue indicates strong structures in multiple directions \rightarrow corners

- Corners: local maxima of the second eigenvalue






- Problem: Detected corners depend on the image scale




- Consider a **Gaussian pyramid** of smoothed images
- The characteristic scale can be computed based on the **Laplacian**:
(Lindeberg 1998)

$$\sigma_c = \operatorname{argmax}_{\sigma} \left(\sigma^2 \cdot |\partial_{xx}(K_{\sigma} * I) + \partial_{yy}(K_{\sigma} * I)| \right)$$

- Yields an estimate of the scale shift between two images
- Uniqueness is not ensured
 - There may be multiple local maxima
 - Even the global maximum need not be unique
- Maximum operator is not robust
 - a little noise can lead to a very different outcome



Authors: Krystian Mikolajczyk and Cordelia Schmid





Authors: Krystian Mikolajczyk and Cordelia Schmid

- Alternative to Harris-Laplace detector
- Considers local maxima of the Laplacian in scale space:

$$x^*, y^*, \sigma^* = \operatorname{argmax}_{x, y, \sigma} \left(\sigma^2 \cdot |\partial_{xx}(K_\sigma * I(x, y)) + \partial_{yy}(K_\sigma * I(x, y))| \right)$$

- Advantage: Does not mix apples and oranges (corner detector and Laplacian)
- Laplacian focuses on blobs rather than corners
 - complementary information
 - one might be interested in using both

- Positive issue of interest points:
 - Significantly reduced complexity
 - With 100 detected points in both images, one has to compare only 10000 patches instead of 96 billion(!) in 640x480 images
- Negative issues:
 - Non-dense displacement fields (important matches might be missed)
 - Corresponding patches can be slightly shifted
- Further problems (independent of interest points)
 - Patches in both images may look very different due to:
 - Rotations
 - Projective transformations (different viewing angles)
 - Lighting changes (shadows, flickering)
 - Blurring
 - Subpixel accuracy not available






Example: projective transformation



Example: lighting changes




Example: blurring



- **Local descriptors:** vectors that contain information about the local neighborhood of an interest point
- Simplest local descriptor: block of a certain size centered at the interest point
- Goal: design local descriptors that are **invariant** under the mentioned transformations


- Use characteristic scale from interest point detection
- Choose and normalize the size of the blocks, such that structures have the same scale in both images



- Affine transformation:


$$f(x) = Ax + t$$

- Maps a circle to an ellipse (or vice-versa):



- Approximation of a projective transformation
- Parameters can be estimated, e.g., from a region detector (maximally stable extremal regions)



- Maximally stable extremal regions
(Matas et al. 2002)
 - Regions encircled by large gradients
 - Obtained by watershed-like algorithm




Maximally stable extremal regions and fitted ellipses. Author: Andrea Vedaldi

- Apart from scale also yields elongation of fitted ellipses
→ allows for affine invariance

- Alternative to a normalized neighborhood:
derive invariant features within the fixed block
- Gray value histogram:
 - Rotational invariance
 - Invariant to blurring
 - Sensitive to lighting changes (bad)
 - Significant loss of information (very bad)
- Histogram of the gradient direction (**orientation histograms**)
 - Invariant to (additive) lighting changes
 - Building block of many successful descriptors





- Very popular local descriptor (several variants exist)
- Based on local assembly of orientation histograms and adaptive local neighborhoods





Author: David Lowe

- Extract SIFT feature points
 - Strongest responses of Laplacian in scale space
→ position and scale
 - Fit quadratic function to obtain subpixel accuracy
 - Create orientation histogram at selected scale
 - Peak of smoothed histogram estimates orientation
 - In case of two peaks, create two feature points
- Estimation of position, scale, and orientation
- Affine invariance can be provided with MSER
 - In object recognition: dense sampling of such points at all positions and all scales, no rotation invariance




Author: David Lowe

1. Compute gradient orientation and magnitude at each pixel
2. Compute orientation indicator at each pixel
 - Create $N \times M \times 8$ array and initialize with zero
 - Quantize the orientation at each pixel (here 8 bins) and add the respective magnitude to the respective entry in the array
3. Local integration \rightarrow orientation histogram
Smooth array with a Gaussian kernel
4. Smooth in orientation direction
(among neighboring channels)



5. Sample feature vectors from the histogram image

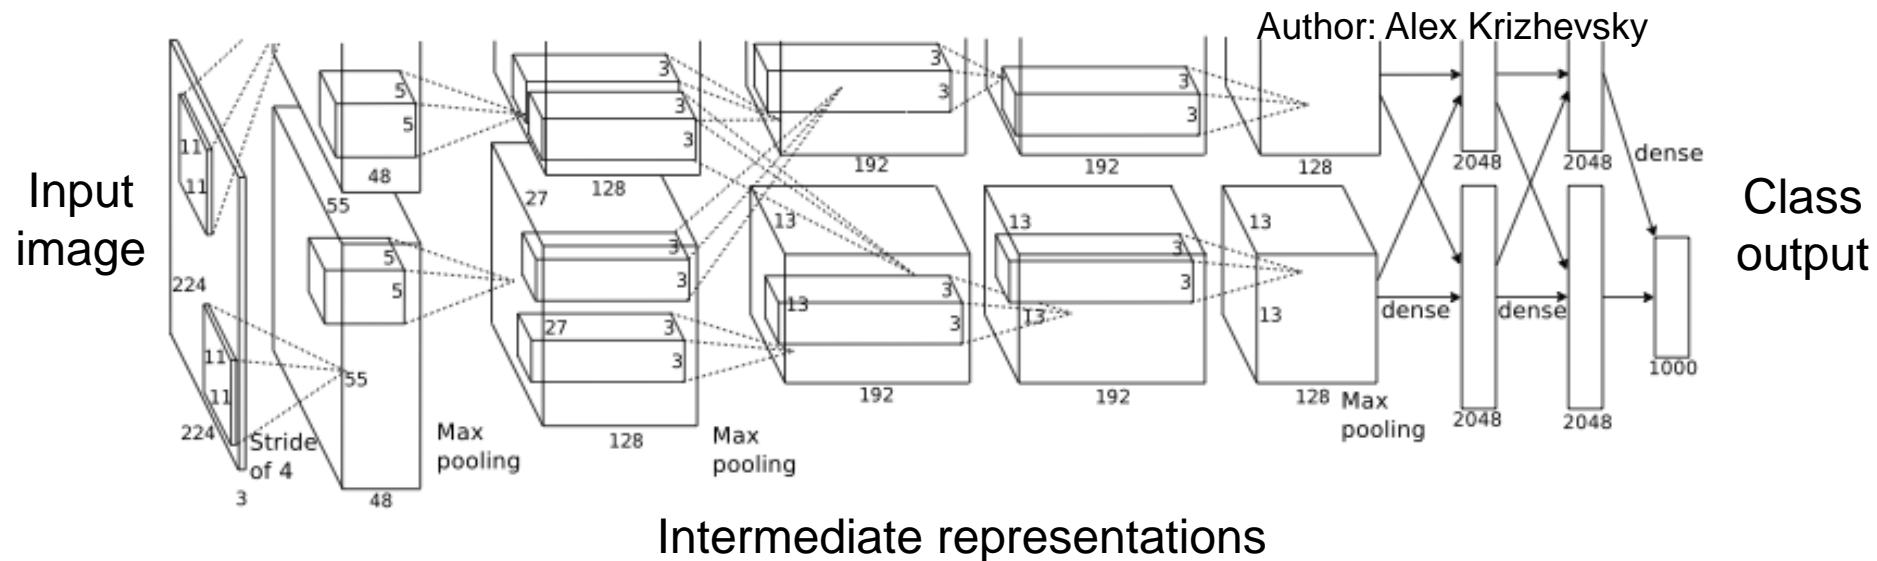
- Original SIFT:
4 pixel spacing, 4x4 histogram array
→ 128-D vector
- HOG for person detection (Dalal-Triggs 2005):
6 pixel spacing, 16x8 histogram array,
9 orientations
→ 1152-D vector



Block consisting of
4 cells


6. Normalize the feature vector

- SIFT: normalize to unit length
- HOG: normalize all **cells** relative to the neighbors of a **block**



- CNNs are trained on large datasets with class labels (e.g. ImageNet with 1M images)
→ network learns a representation that is good for object classification
- Intermediate layer outputs turn out to be good generic descriptors


- Train CNN to discriminate **surrogate classes** (Dosovitskiy et al. 2015)




Seed patch and transformed versions of it make up a surrogate class

- Applied transformations:
translation, rotation, scaling, color, contrast, brightness, blur
- Transformations define invariance properties of the features to be learned by the network

Exemplar-CNN-blur vs SIFT

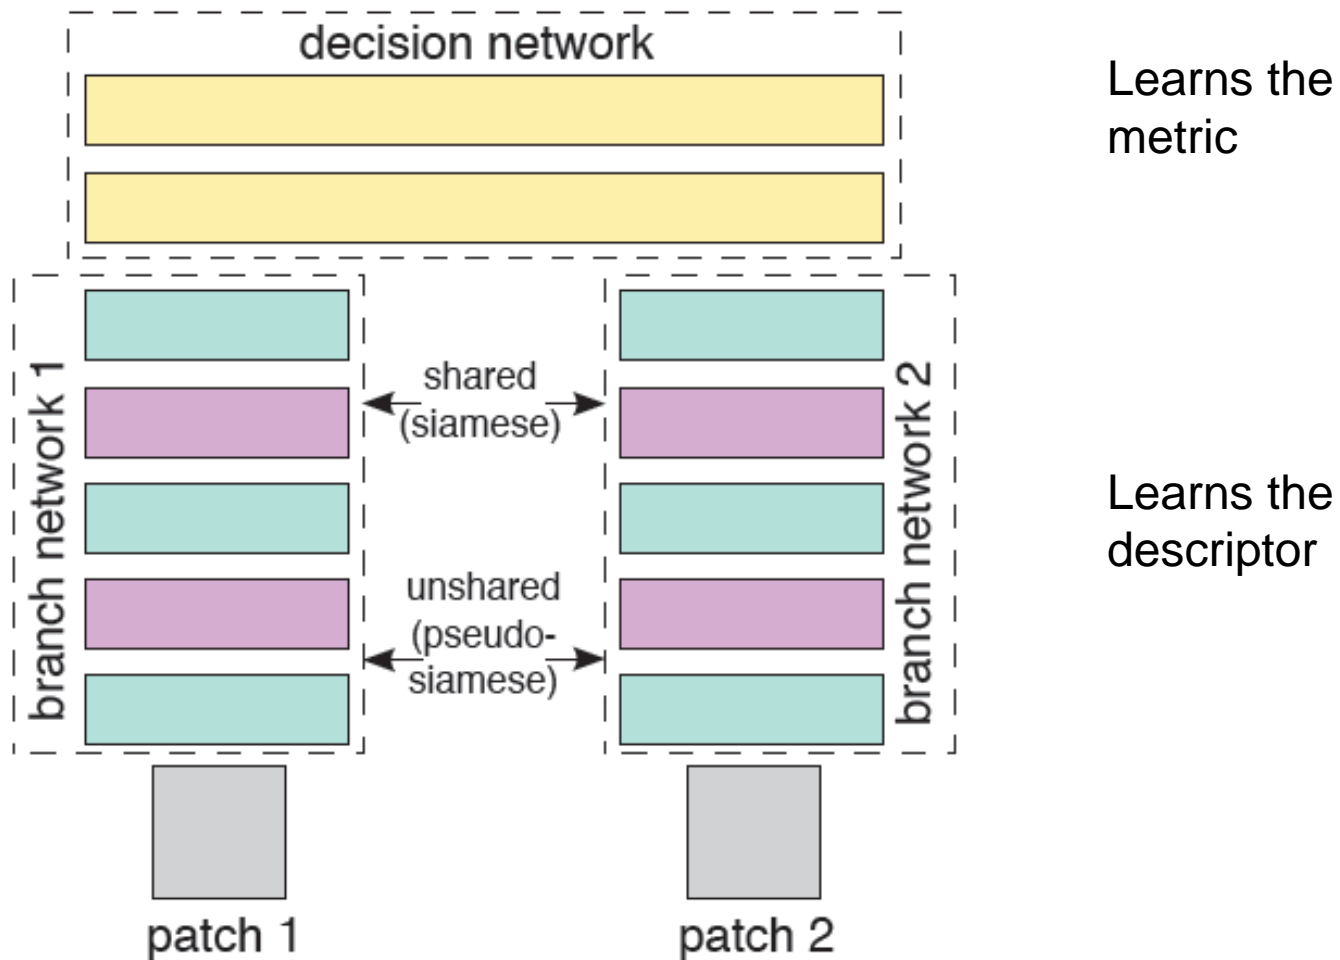


Exemplar-CNN-blur vs AlexNet



Siamese networks

- Trained directly on matching and non-matching patches



Example from Zagoruyko&Komodakis 2015

- Interest points are distinctive points in an image with a significant information content in their neighborhood
- Interest point detection can help establish invariance to certain image transformations.
- Local descriptors describe a local area in the image for the purpose of matching.
- The SIFT descriptor is based on a grid of orientation
- Intermediate layers of ConvNets yield good descriptors

- T. Lindeberg: Feature detection with automatic scale selection, *International Journal of Computer Vision*, 30(2): 79-116, 1998.
- D. G. Lowe: Distinctive image features from scale-invariant keypoints, *International Journal of Computer Vision* 60(2):91-110, 2004.
- J. Matas, O. Chum, M. Urban, T. Pajdla: Robust wide baseline stereo from maximally stable extremal regions, *Proc. British Machine Vision Conference*, 2002.
- N. Dalal, B. Triggs: Histograms of oriented gradients for human detection, *CVPR*, 2005.
- S. Belongie, J. Malik, J. Puzicha: Shape matching and recognition using shape contexts, *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 24(4):509-522, 2002.
- A. Dosovitskiy, P. Fischer, T. Springenberg, M. Riedmiller, T. Brox: Discriminative unsupervised feature learning with convolutional neural networks, *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2016.
- S. Zagoruyko, N. Komodakis: Learning to Compare Image Patches via Convolutional Neural Networks, *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2015.

- Implement the corner detector based on the second eigenvalue of the structure tensor
 - For computing derivatives and for smoothing images you can make use of the predefined filter masks as well as the convolution operations in `CFilter.h`.
 - The structure tensor of a color image is the sum of tensors over all channels
 - See an online math lecture if you do not remember how to compute the eigenvalues of a matrix <http://www.khanacademy.org/>
- Apply the corner detector to the images in `ImageProcessing08Ex03.zip` and play with the parameters
- Implement the dense SIFT/HOG descriptor (without the detector). Use a 4 pixel spacing and a 3x3 grid of histograms. You can ignore scale and rotation invariance and even skip normalization for this exercise.
- Run your corner detector on `tennis500.ppm` and manually select among the interest points the 10 visually most interesting ones. Compute SIFT descriptors for these points.
- Compute SIFT descriptors for all points in `tennis505.ppm`. For each descriptor in `tennis500.ppm` find the best match in `tennis505.ppm` and visualize them in your result image.
- Play with the amount of smoothing, the spacing, and the number of histograms when computing the descriptors.