

# *Image Processing and Computer Graphics*

# *Transparency and Reflection*

Matthias Teschner

Computer Science Department  
University of Freiburg

Albert-Ludwigs-Universität Freiburg

# Outline


---

- transparency
- reflection

# Introduction

---

- simplified transparency model
  - semitransparent objects are filters / attenuators of occluded objects
  - refraction and object thickness are neglected
- algorithms are based on
  - stipple patterns
  - color blending per pixel



[Cass Everitt: Interactive Order-Independent Transparency]

# Stipple Patterns

- screen-door transparency
- transparent object is rendered with a fill / stipple pattern, e.g. checkerboard (pattern of opaque and transparent fragments)
- limited number of fill patterns results in limited number of transparency levels
- aliasing artifacts
- simple method


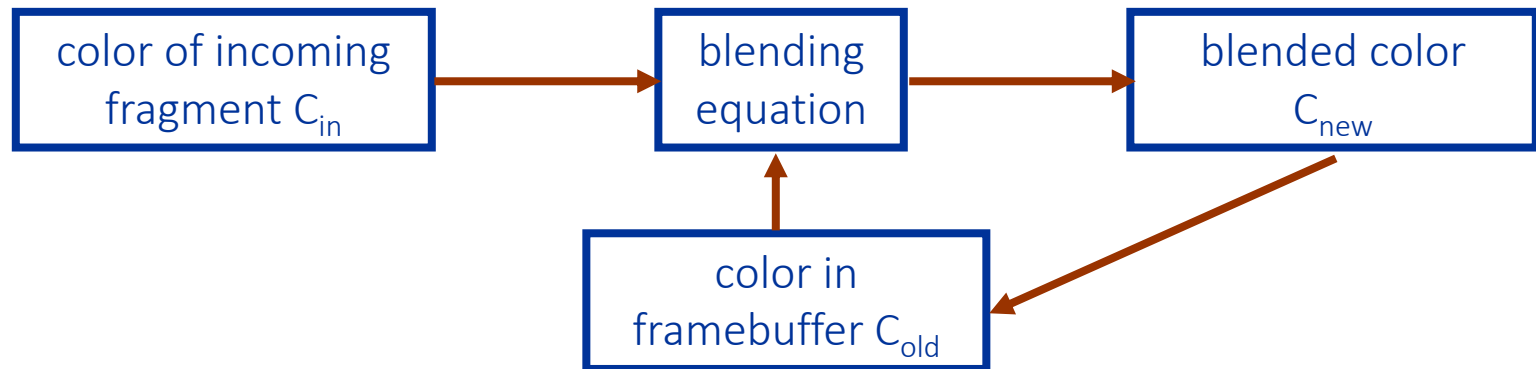


illustration of a stipple pattern

# Color Blending

- combine fragment color with the framebuffer content



- color  $C_{old}$  is replaced by  $C_{new}$
- blending equation:  $C_{new} = \alpha_{in} \cdot C_{in} + \alpha_{old} \cdot C_{old}$


# Color Blending

---


- alpha value
  - describes the opacity of a fragment,  
1 - opaque, 0 - transparent
  - stored together with RGB color in a 4D vector (RGBA)
- blending equation for transparency
  - $\mathbf{C}_{\text{new}} = \alpha_{\text{in}} \cdot \mathbf{C}_{\text{in}} + (1 - \alpha_{\text{in}}) \cdot \mathbf{C}_{\text{old}}$
  - over operator
  - $\alpha_{\text{in}} = 0$  -  $\mathbf{C}_{\text{old}}$  is not changed
  - $\alpha_{\text{in}} = 1$  -  $\mathbf{C}_{\text{old}}$  is replaced by  $\mathbf{C}_{\text{in}}$
  - $0 < \alpha_{\text{in}} < 1$  -  $\mathbf{C}_{\text{old}}$  is replaced by a mix of  $\mathbf{C}_{\text{in}}$  and  $\mathbf{C}_{\text{old}}$
  - only the alpha value of the incoming fragment matters

# Color Blending


- order matters




B over A



A over B



A over B over C



B over A over C

# Outline

---


- transparency
  - depth ordering
  - binary space partitioning
  - depth peeling
- reflection



# Depth Ordering

---

- polygons / fragments have to be rendered in sorted depth order
- hardware generally renders in object order
- depth test only returns the nearest fragment per pixel, sorting is not realized
- intersecting polygons have to be handled
- dynamic scenes require re-sorting



[Cass Everitt: Interactive Order-Independent Transparency]

# *Depth Ordering for Convex Objects*

---

- exactly two depth layers for arbitrary viewing directions
- first depth layer defined by front faces
- second depth layer defined by back faces
- algorithm
  - render back faces in a first pass
  - blend with front faces in a second pass

# *Depth Ordering for Arbitrarily Shaped Objects*

---

- object-space methods
  - use pre-computed spatial data structures
  - e.g., binary space partition tree (BSP tree)
  - useful for static geometry
  - varying viewer positions and orientations can be handled
- screen-space methods
  - employ the functionality of the rendering pipeline
  - several rendering passes compute depth layers
  - final pass renders the ordered depth layers
  - useful for dynamic / deforming geometry and arbitrary views
  - no pre-computation is required / can be employed


# Outline

---

- transparency
  - depth ordering
  - binary space partitioning
  - depth peeling
- reflection


# Binary Space Partitioning BSP

- BSP tree is a hierarchical spatial data structure
- 3D space is subdivided by means of arbitrarily oriented planes
- nodes represent planes
- leaves represent convex space cells
- applications
  - visible surface algorithm
  - depth sorting
  - collision detection




# Generation of the BSP Tree

- the BSP tree is pre-computed for static scenes
- all planar primitives are represented in the tree
- balancing is less important, as the entire tree has to be queried (all primitives are rendered)





# Query of the BSP Tree

- motivation
  - a viewer is on the near side of a plane
  - a polygon on the far side of this plane cannot occlude the plane or any polygon on the near side
- back to front rendering
  - render far branch of the viewpoint
  - render root (node) polygon
  - render near branch of the viewpoint
  - recursively applied to sub-trees



# Query of the BSP Tree

- back to front rendering
- viewpoint is in 1-
- rendering of 1+, 1, 1-
- rule recursively applied to 1+ and 1-
- viewpoint is in 3+
  - rendering of 3, 2b
- viewpoint is in 4-
  - rendering of 2a, 4





# *BSP Tree - Discussion*

---

- not only visible surface generation, but depth sorting of all primitives per pixel position
- additional data structure
- can be pre-computed
- requires polygon splits
- dynamic scenes require an update of the data structure

# Outline

---

- transparency
  - depth ordering
  - binary space partitioning
  - depth peeling
- reflection


# Concept

---

- motivation
  - use the functionality of the rendering pipeline
  - several rendering passes compute depth layers
  - final pass renders the ordered depth layers
  - useful for dynamic / deforming geometry
  - no pre-computation is required / can be employed
- algorithm
  - first render pass gives the front-most fragment color / depth
  - each successive render pass extracts the fragment (with color and depth) for the next-nearest fragment on a per pixel basis (screen-space approach)
  - two depth buffers are used

# Concept

- object is rendered once for each depth layer
  - depth complexity is the max number of layers per pixel position
- two separate depth tests per fragment
  - must be farther than the one in the previous layer ( $d_1$ )
  - must be the nearest of all remaining fragments ( $d_2, d_3, d_4$ )




[Bruno Heidelberger]

# Depth Layers - 2D

- depth peeling strips away one depth layer with each successive rendering pass
- illustration
  - bold black lines - frontmost (leftmost) surfaces
  - thin black lines – hidden surfaces
  - light grey lines – “peeled away” surfaces


[Cass Everitt:  
Interactive  
Order-Independent  
Transparency]




# Depth Layers - 3D




depth layer 1



depth layer 2



depth layer 3



depth layer 4

[Cass Everitt: Interactive Order-Independent Transparency]

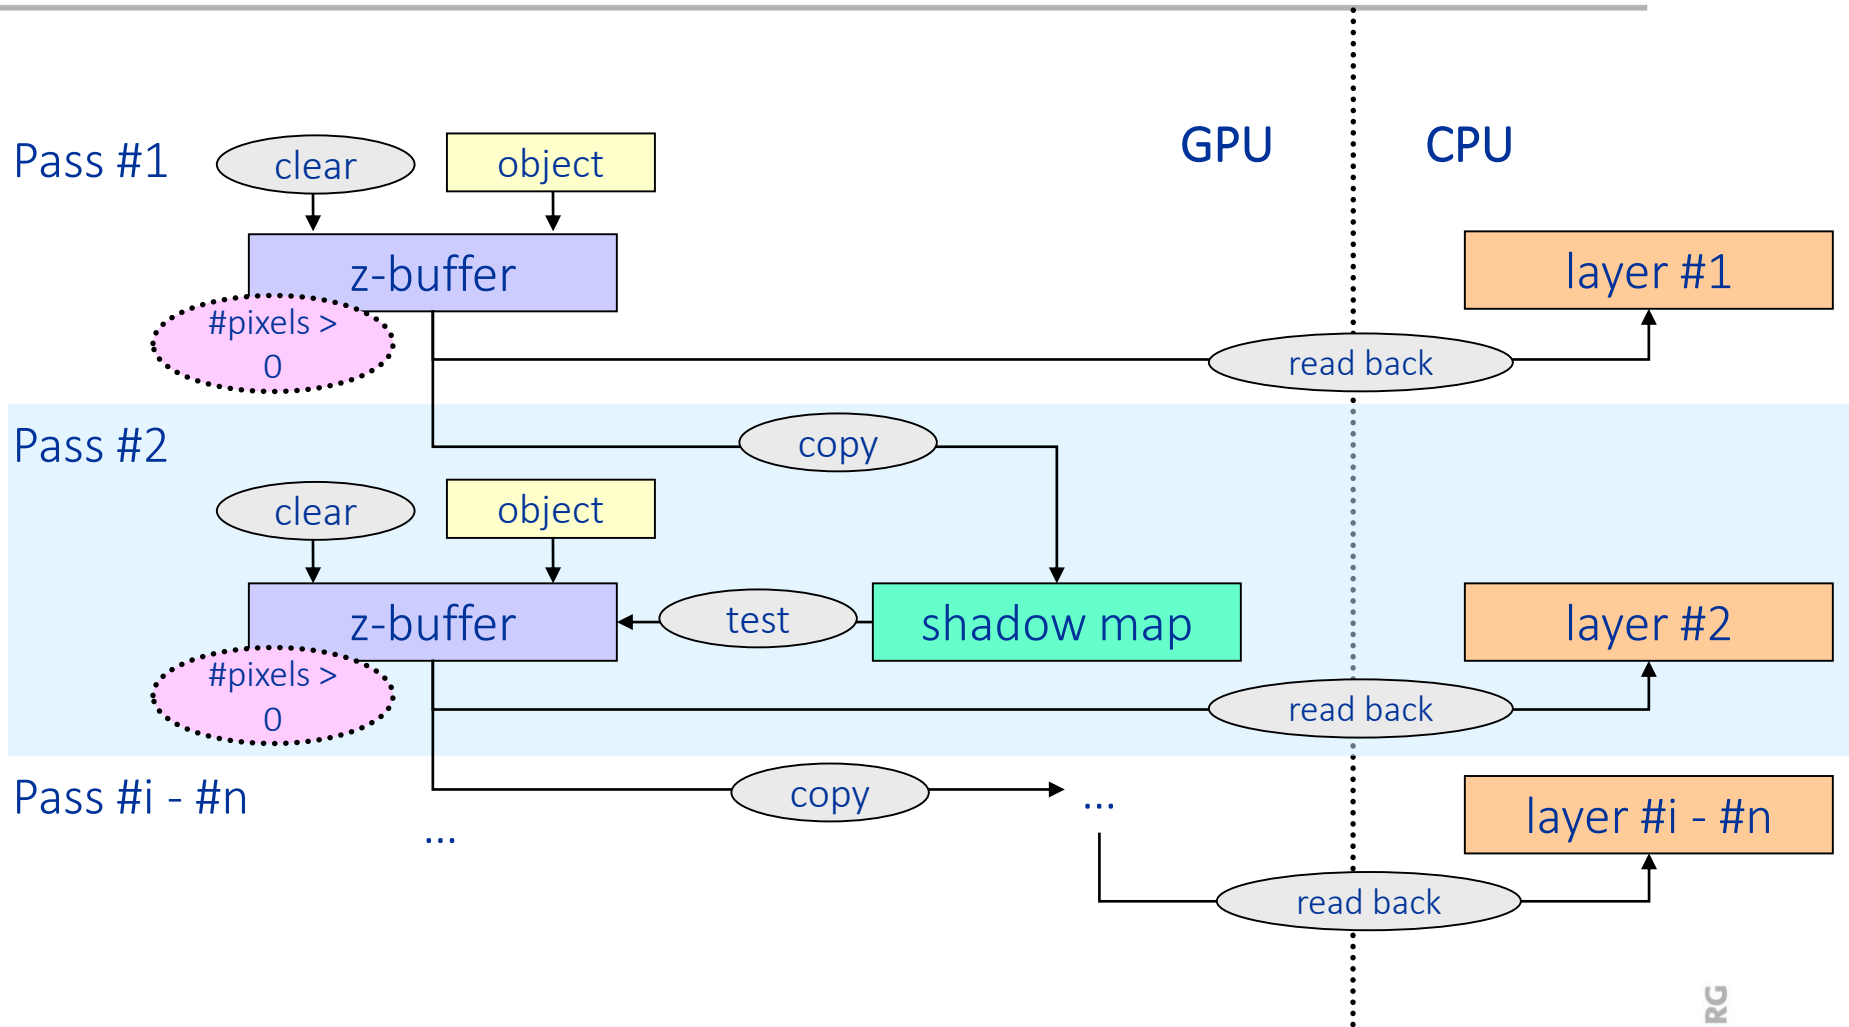
# *Implementation*

## *Based on Shadow Mapping*

---

- two depth tests (depth buffers ) are required
- e.g., shadow mapping can be used to realize a second depth buffer
- in contrast to the depth buffer, the shadow map
  - is not tied to the camera position
  - is not writeable during depth test
  - does not discard fragments
- with respect to depth peeling
  - the shadow map is tied to the camera position
  - copy functionality to depth buffer is employed

# Implementation Based on Shadow Mapping




[Bruno Heidelberger]




# Results


- quality and performance are determined by the number of generated depth layers




one  
depth  
layer



two  
depth  
layers



three  
depth  
layers



four  
depth  
layers

[Cass Everitt: Interactive Order-Independent Transparency]

# *Depth Peeling - Discussion*

---

- screen-space algorithm
- multiple rendering passes generate depth layers per pixel position
- view dependent (in contrast to the BSP approach)
- appropriate for dynamic scenes
- quality and performance are determined by the number of rendering passes (in the discussed implementation)

# *Transparency - Summary*

---

- simplified transparency model
- algorithms based on
  - stipple patterns
  - color blending
- for blending, depth-sorted primitives are required
- BSP tree
  - object space algorithm with one rendering pass
  - appropriate for static scenes
- depth peeling
  - screen space algorithm with multiple rendering passes
  - appropriate for dynamic scenes


# Outline

---

- transparency
- reflection
  - planar surfaces
  - arbitrary surfaces

# Law of Reflection

- angle of incidence is equal to the angle of reflection



[Akenine-Moeller et al.: Real-time Rendering]

# Generation of Reflected Geometry

- original and reflected geometry is rendered
- reflected geometry is generated with respect to the reflection plane with surface normal  $\mathbf{n} = (n_x, n_y, n_z)$  and a point  $\mathbf{p}$  on the reflection plane


$$\mathbf{M}_{(\mathbf{n}, \mathbf{p})} = \begin{pmatrix} 1 - 2n_x^2 & -2n_x n_y & -2n_x n_z & 2n_x(\mathbf{n} \cdot \mathbf{p}) \\ -2n_x n_y & 1 - 2n_y^2 & -2n_y n_z & 2n_y(\mathbf{n} \cdot \mathbf{p}) \\ -2n_x n_z & -2n_y n_z & 1 - 2n_z^2 & 2n_z(\mathbf{n} \cdot \mathbf{p}) \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

e.g.


$$\mathbf{M}_{((0,1,0),(0,0,0))} = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & -1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

# Implementation

- render reflected geometry with reflected illumination
- render semi-transparent reflection plane, e.g. with color blending
- render original geometry
- render reflection plane to stencil
- render reflected geometry where stencil is set
- ...



reflection  
rendering  
without  
stenciling



reflection  
rendering  
with  
stenciling

[Akenine-Moeller et al.: Real-time Rendering]

# Outline


---

- transparency
- reflection
  - planar surfaces
  - arbitrary surfaces




# Environment Mapping



- e.g. cube mapping
- approximates reflections of the environment on arbitrary surfaces



place a viewer in a scene



generate the environment texture from six view directions




apply the texture to an object at the position of the viewer

[Wikipedia: Cube Mapping]

# Environment Mapping

- environment is projected onto an object-embedding shape, e. g. sphere or cube
- view-dependent mapping
  - dependent on viewing and reflection direction
- approximate implementation of reflections off arbitrary surfaces




[Rosalee Wolfe]

# Environment Mapping

## Motivation


---

- the Phong illumination model (a local model) does not take into account reflections
- Raytracing (a global model) traces rays off the object into the world to obtain reflections



# Environment Mapping

- environment mapping approximates this process by capturing the environment in a texture map and using the reflection vector to index into this map
- cannot handle changing reflections of moving objects




# *Environment Mapping - Steps*

---

- generate or load a 2D map of the environment
- for each fragment of a reflective object, compute the normal  $\mathbf{n}$
- compute the reflection vector  $\mathbf{r}$  from the view vector  $\mathbf{v}$  and the normal  $\mathbf{n}$  at the surface point
- use the reflection vector to compute an index into the environment map that represents the objects in the reflection direction
- use the texel data (texture value) from the environment map to color the current fragment

# Cubic Environment Mapping


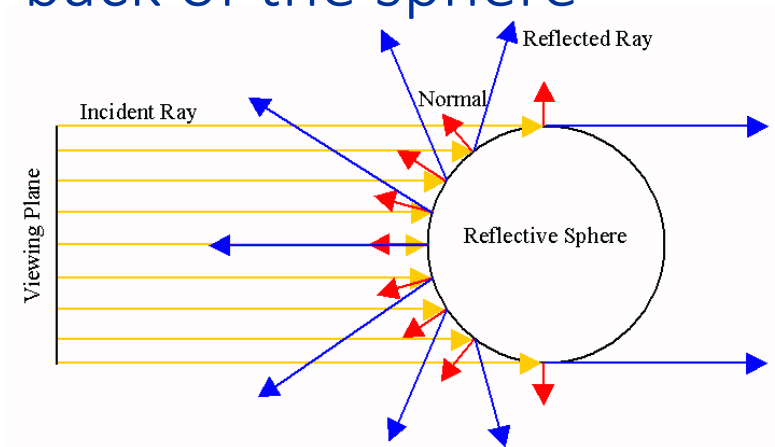
- the map is constructed by placing a camera at the center of the object and taking pictures in 6 directions



[Mizutani, Reindel: Environment Mapping Algorithms, Reindel Software]

# Spherical Environment Mapping

- the map is obtained by orthographically projecting an image of a mirrored sphere
- map stores colors seen by reflected rays
- sphere map contains information about both, the environment in front of the sphere and in back of the sphere



[Mizutani, Reindel: Environment Mapping Algorithms, Reindel Software]

# *Spherical Environment Mapping*

---

- the map can be obtained from a synthetic scene by
  - Raytracing
  - warping automatically generated cubic maps
- the map can be obtained from the real world by
  - photographing an actual mirrored sphere



[<http://www.oakcorp.net/chaos/hdri.shtml>]



# *Spherical Environment Mapping*

---



[Mizutani, Reindel: Environment Mapping Algorithms, Reindel Software]

# Spherical Environment Mapping

---

- to map the reflection vector to the sphere map, the following equations are used based on the reflection vector  $\mathbf{r}$
- in contrast to cube mapping, a generalized equation can be used

$$\begin{pmatrix} u \\ v \end{pmatrix} = \begin{pmatrix} \frac{r_x}{2\sqrt{r_x^2 + r_y^2 + (r_z + 1)^2}} + \frac{1}{2} \\ \frac{r_y}{2\sqrt{r_x^2 + r_y^2 + (r_z + 1)^2}} + \frac{1}{2} \end{pmatrix}$$

# *Spherical Environment Mapping*

---

- disadvantages
  - maps are hard to create on the fly
  - sampling is non-linear, non-uniform
  - sampling is view-point dependent
- advantages
  - no interpolation across map seams

# *Environment Mapping*

## *Discussion*

---

- object should be small compared to the environment
- issues with self-reflections and non-convex objects
- separate map for each object
- maps may need to be changed in case of a changing viewpoint due to non-uniform sampling
- translated objects might require a map update

# Reflection - Summary

---

- planar reflections
  - generation of reflected geometry with reflected lighting
  - rendering of reflected geometry, reflection plane, original geometry
  - blending and stenciling is employed
- arbitrarily shaped reflectors
  - approximate reflections with environment mapping
  - cube maps
  - sphere maps
  - works best for distant environments without translation of objects
  - issues with sampling and concave objects